Deep Neural Networks for Information Extraction

Tomáš Gogár, Petr Baudiš

Information Extraction

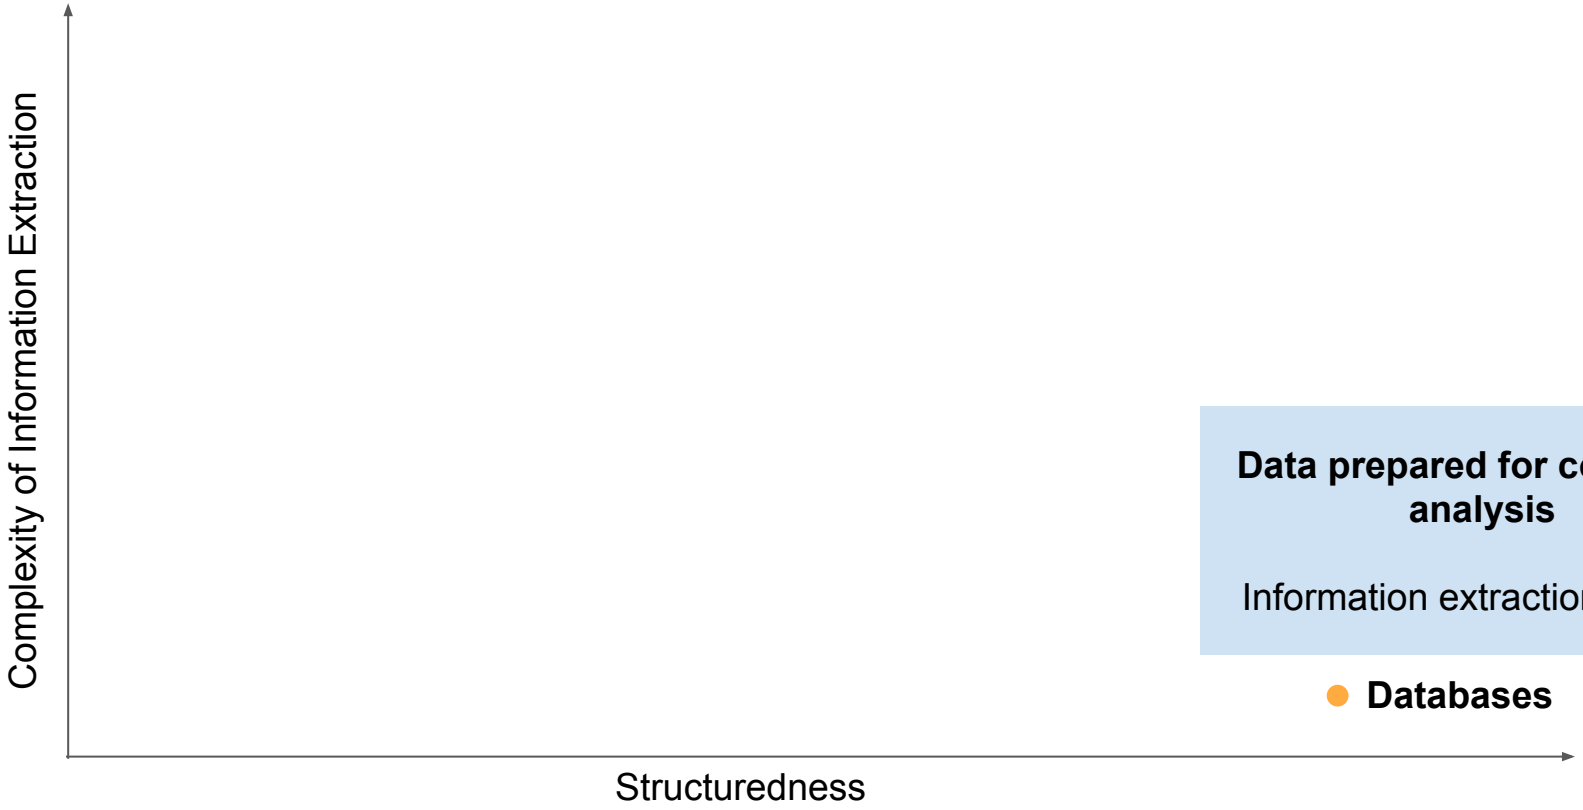

Information Extraction

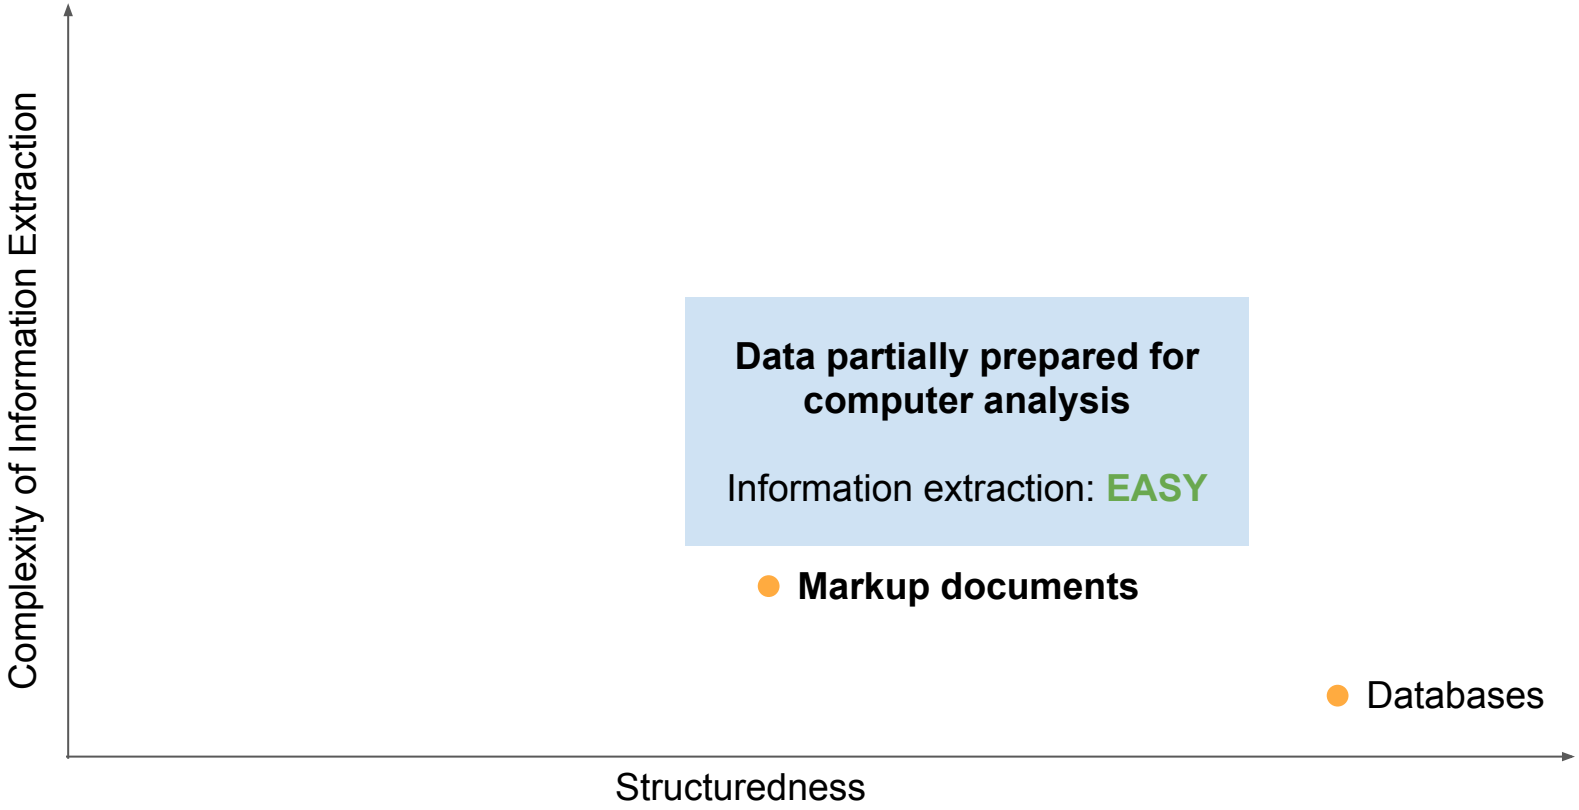

Information Extraction

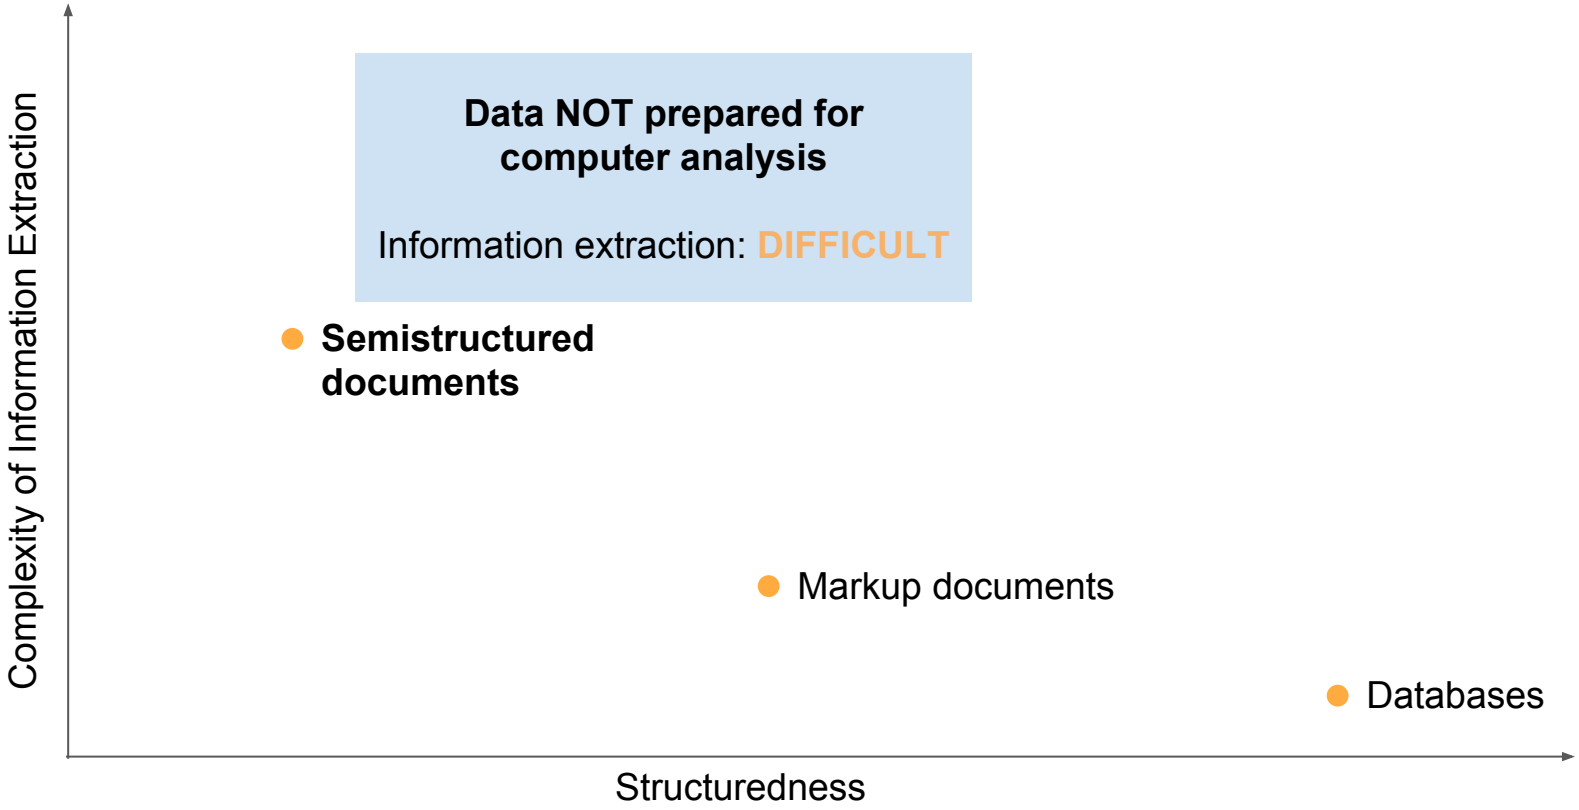

Information Extraction

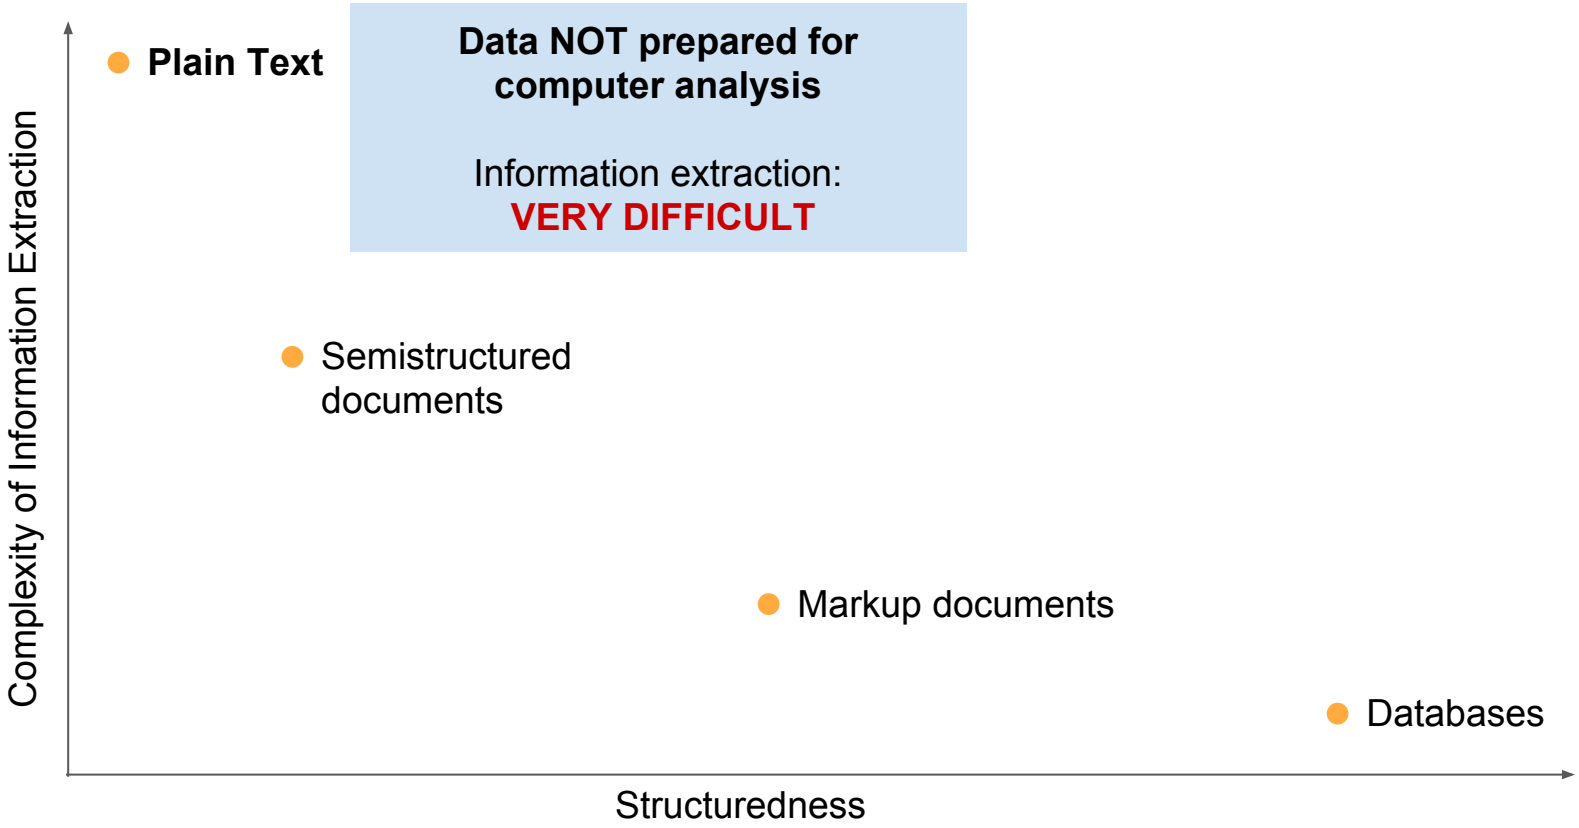

Information Extraction

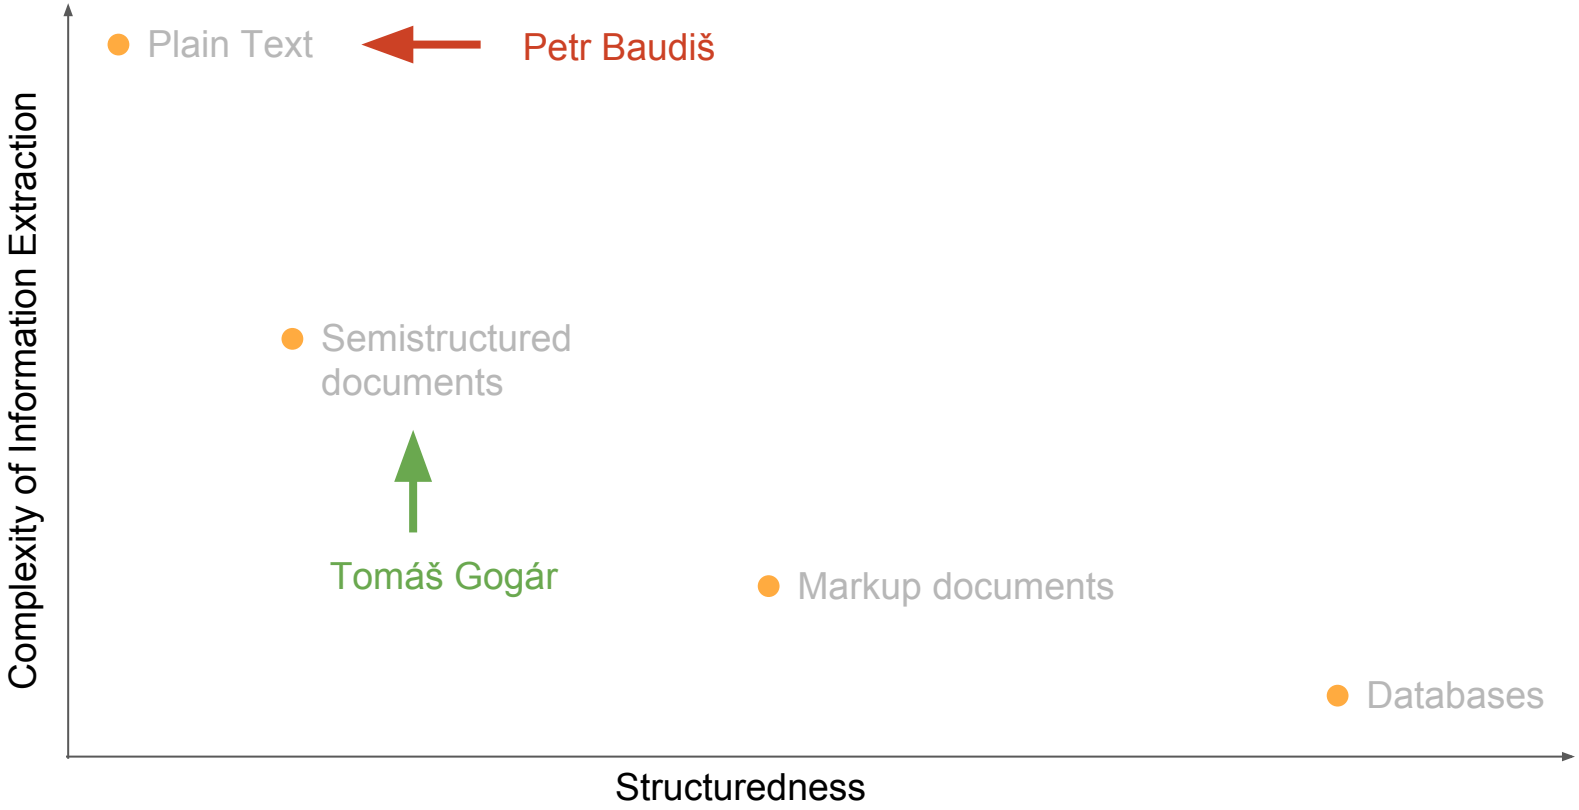

Web Information Extraction - Current systems

- Web pages are created from Templates
- **Learn template structure ⇨ Extract Information**

Web Information Extraction - Current systems

- Web pages are created from Templates
- **Learn template structure ⇨ Extract Information**
- **Template learning:**
  - **Manual annotation - Scraping**

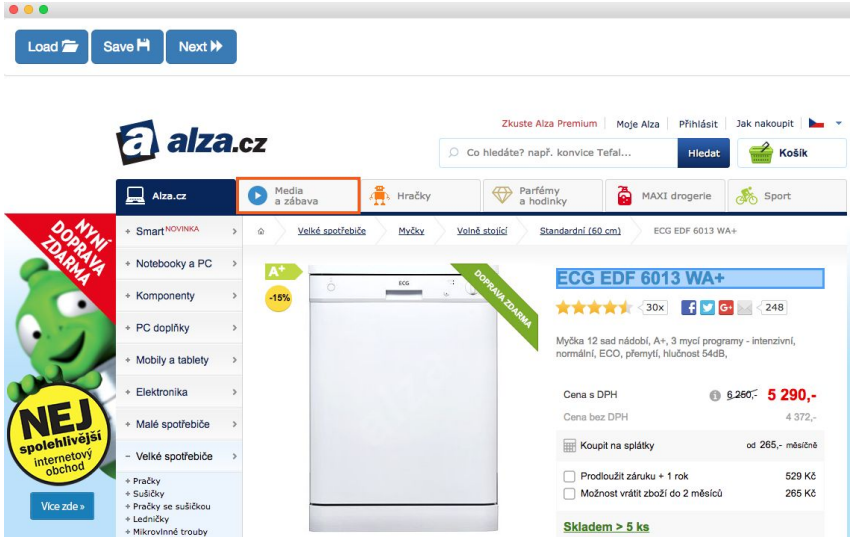

Web Information Extraction - Current systems

- Web pages are created from Templates
- **Learn template structure ⇨ Extract Information**
- **Template learning:**
  - **Manual annotation**
  - **Automatic learning - repeated patterns**

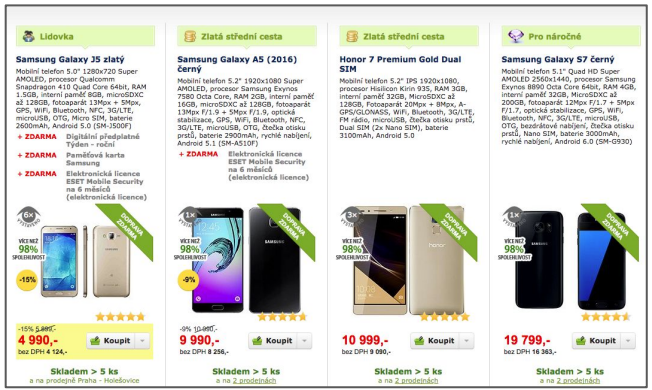

In a page

Across website

Web Information Extraction - Current systems

- Web pages are created from Templates
- **Learn template structure ⇨ Extract Information**
- **Template learning:**
  - **Manual annotation**
  - **Automatic learning - repeated patterns**

It’s just a hack!

What matters in Information Extraction

What is written?

Where it is written?

How it is written?

What matters in Information Extraction

What is written?

Where it is written?

How it is written? **[Screenshot]**

What matters in Information Extraction

What is written?

Where it is written?

} **[SPATIAL BAG-OF-WORDS]**

How it is written? **[Screenshot]**

Intro: Bag-of-Words

Text representation often used in NLP:

(1) John likes to watch movies. Mary likes movies too.

(2) John also likes to watch football games.

Vocabulary:

"John"	1
"likes"	2
"to"	3
"watch"	4
"movies"	5
"also"	6
"football"	7
"games"	8
"Mary"	9
"too"	10

Vectors:

Document1 = [1, 2, 1, 1, 2, 0, 0, 0, 1, 1]

Document2 = [1, 1, 1, 1, 0, 1, 1, 1, 0, 0]

Intro: Hashing Trick

Text representation often used in NLP:

(1) John likes to watch movies. Mary likes movies too.

(2) John also likes to watch football games.

~~**Vocabulary:**~~

Hashing function

h(john**) = 3**
h(likes**) = 1**
…

- Does not need vocabulary **:-)**
- Arbitrary size of result vector **:-)**
- Collisions **:-(**

Spatial Bag-of-Words

- We do not process text as a whole
- We process each TEXT NODE individually

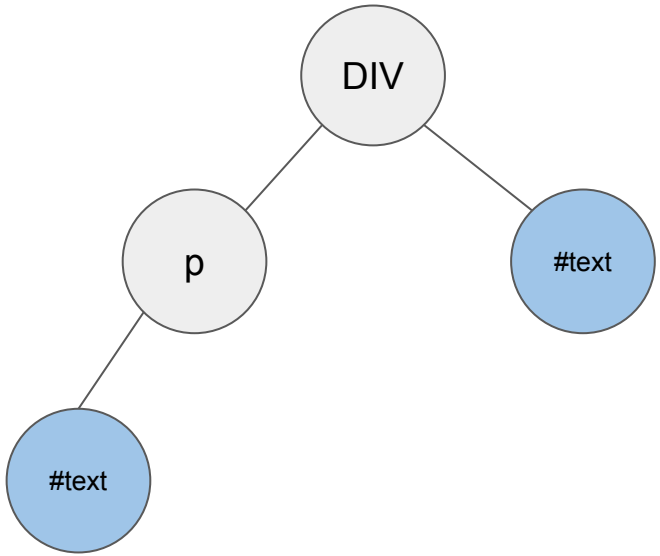

Spatial Bag-of-Words

- We do not process text as a whole
- We process each TEXT NODE individually

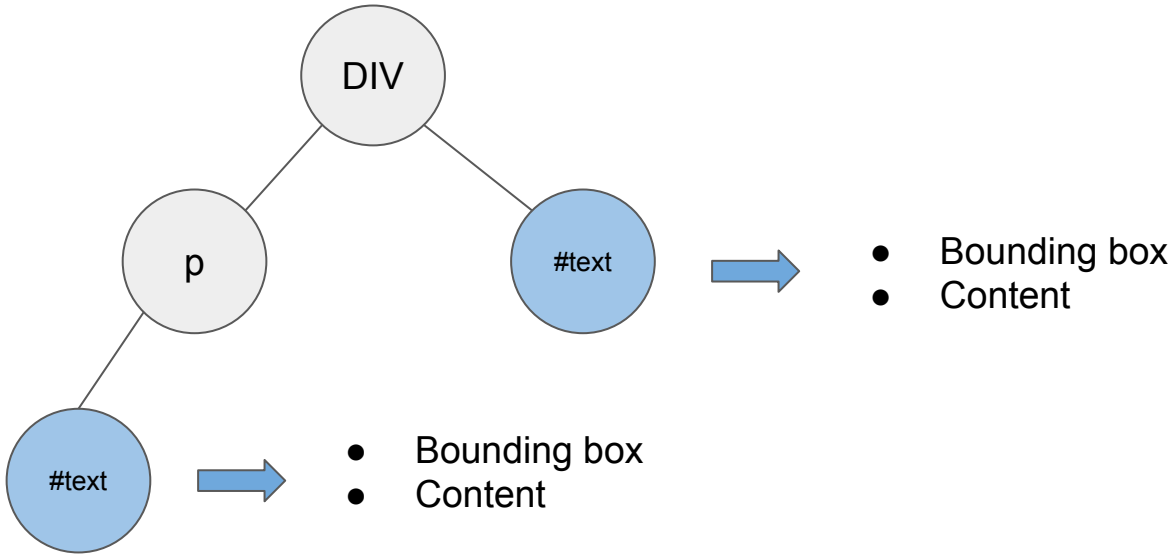

Spatial Bag-of-Words

- We do not process text as a whole
- We process each TEXT NODE individually

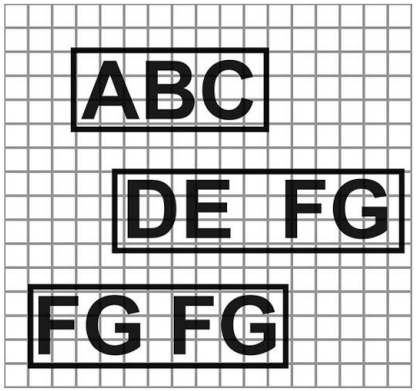

Spatial Bag-of-Words

- We do not process text as a whole
- We process each TEXT NODE individually

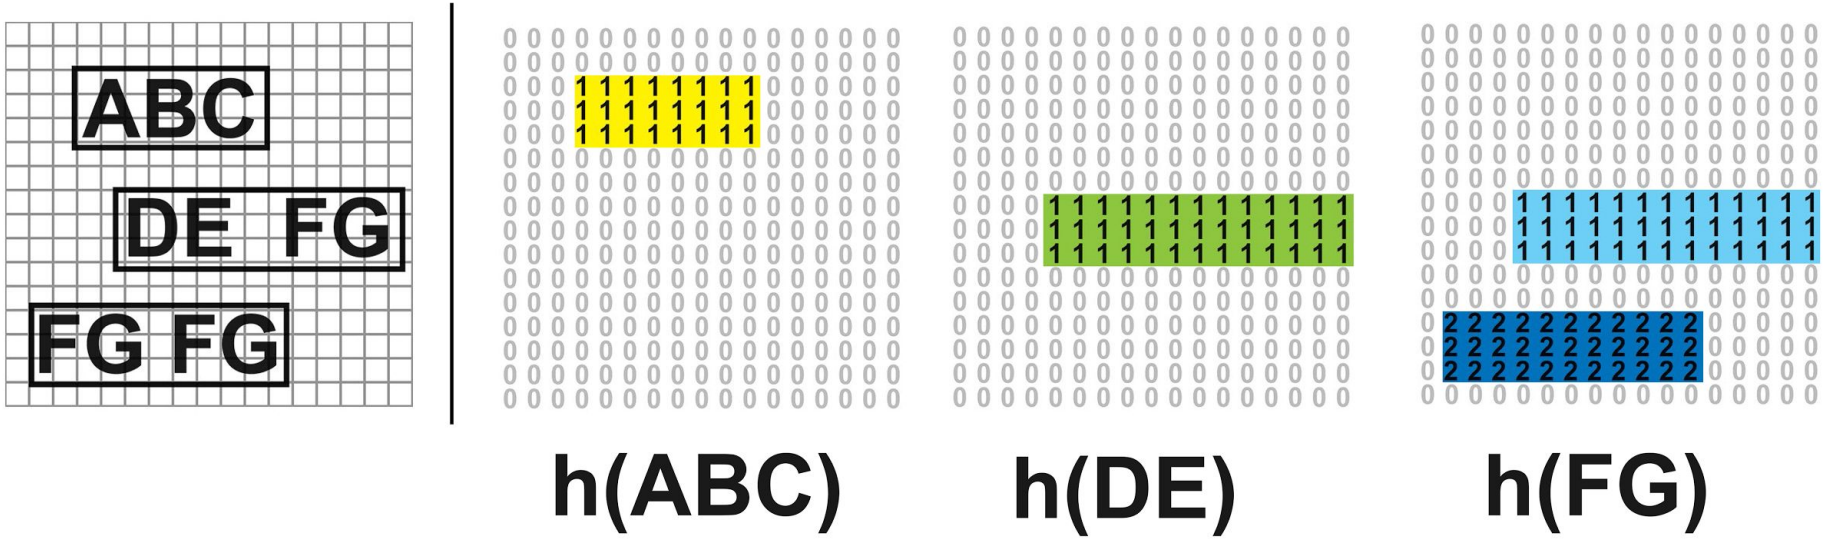

Spatial Bag-of-Words

- We do not process text as a whole
- We process each TEXT NODE individually

TEXT ENCODED IN TENSOR

(SAME AS IMAGE)

Net architecture

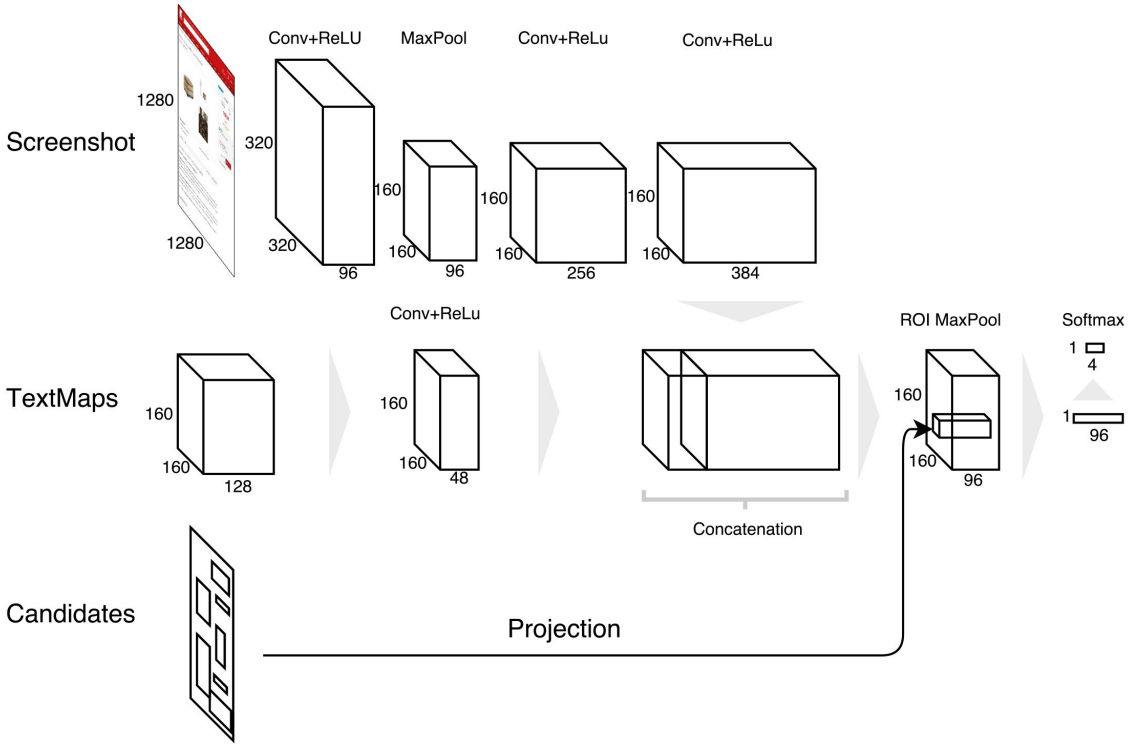

Problem: Uses only local information!

Spatial likelihood

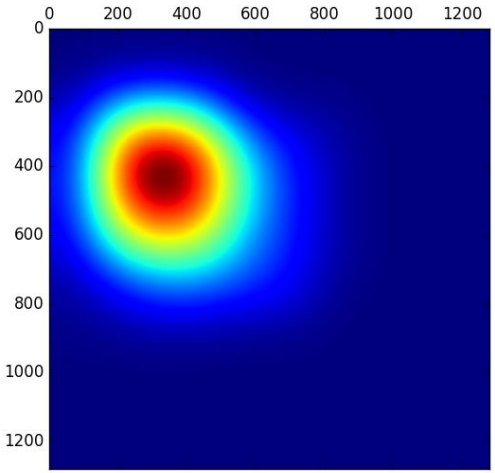

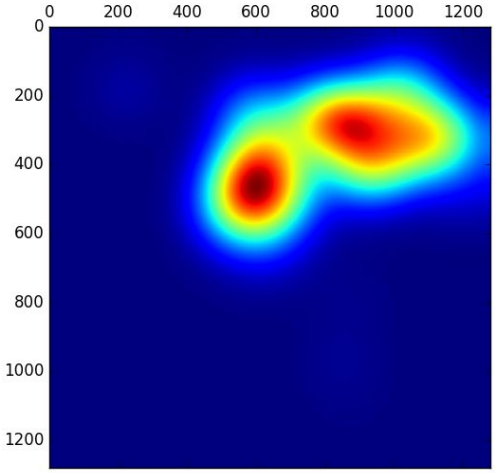

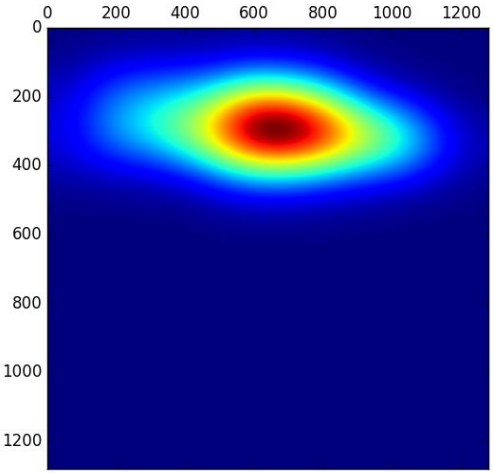

Final system

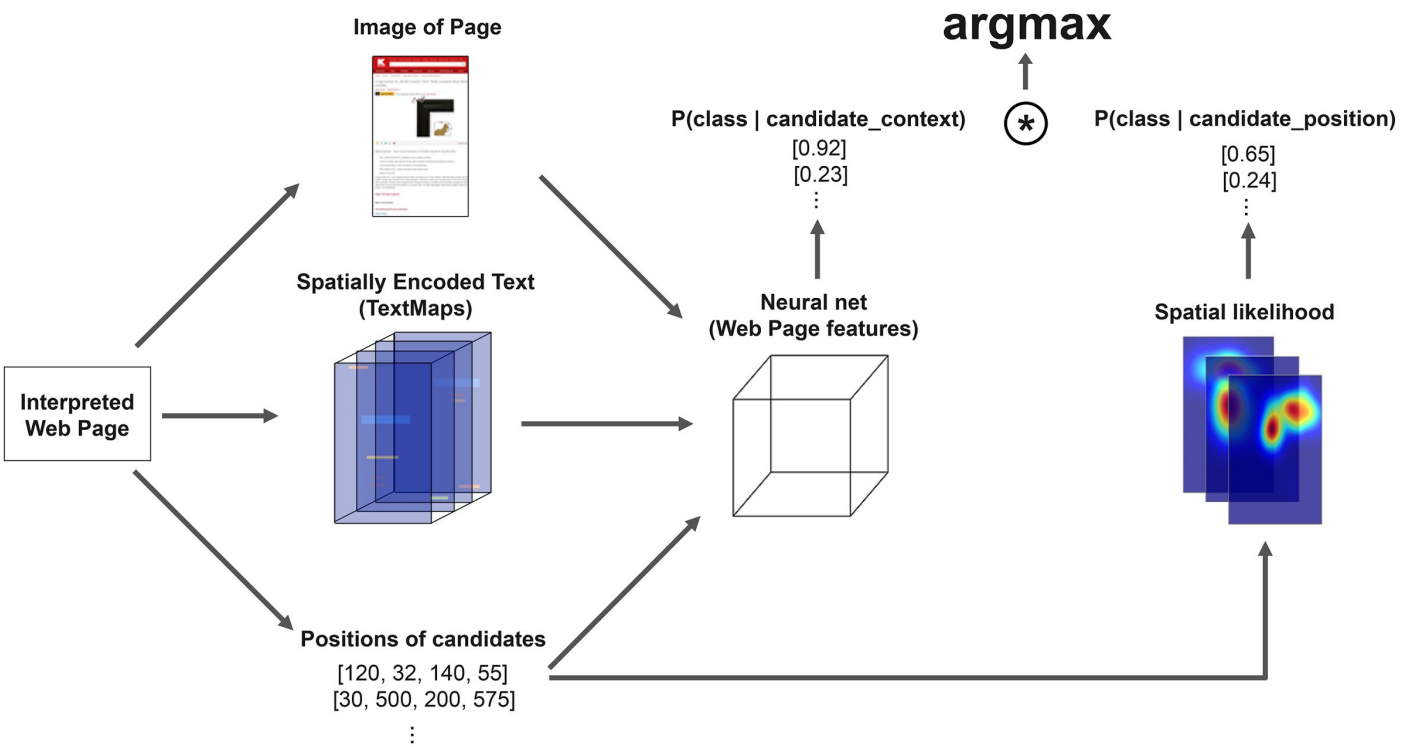

Results

Algorithm	Image Accuracy	Price Accuracy	Name accuracy
NeuralNet+Spatial.	**98.7±1.6**	**95.3±6.6**	**87.1±15.0**
NeuralNet	95.9±2.9	86.2±9.3	78.4±19.0
Baseline: Heuristic + Spatial.	63.7±20.1	73.6±18.8	34.4±20.5
Baseline: Spatial	46.5±18.7	9.7±14.4	12.2±12.0

Table 3. Comparison of algorithms: mean and standard deviation of accuracy across 10 splits (in percents).

Neural net inputs	Image Accuracy	Price Accuracy	Name accuracy
Screenshot + TextMap	95.9±2.9	86.2±9.3	78.4±19.0
Screenshot	93.5±7.4	73.3±19.4	73.4±16.0
TextMap	41.4±18.6	77.0±17.9	49.4±18.0

Table 4. Neural Net with different input data: mean and standard deviation of accuracy across 10 splits (in percents).

Results

Fig. 5. Examples of *current price* detection.

Results

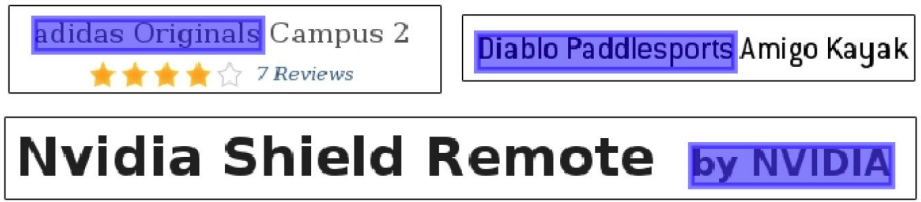

Fig. 6. Examples of product names divided into two parts (manufacturer + model).

Future work

Machine learning:

- Solve global position problem: Attention network?
- Try to learn text features
- Try other similar tasks: ex. classification

Practical problems:

- Popup windows
- Information distributed in multiple DOM-Elements

Source code:

github.com/gogartom/TextMaps